

A PROBABILISTIC GENERATIVE METHOD FOR SAFE PHYSICAL SYSTEM CONTROL PROBLEMS

Peiyan Hu^{2§*} Xiaowei Qian^{1§*} Wenhao Deng¹ Rui Wang^{3§} Haodong Feng¹
 Ruiqi Feng¹ Tao Zhang¹ Long Wei¹ Yue Wang⁴ Zhi-Ming Ma² Tailin Wu^{1†}

¹School of Engineering, Westlake University,

²Academy of Mathematics and Systems Science, Chinese Academy of Sciences,

³Fudan University

⁴Microsoft AI4Science

{hupeiyan, wutailin}@westlake.edu.cn, {xiaoweiqian0311}@gmail.com

ABSTRACT

Controlling complex physical systems is a crucial task in science and engineering, often requiring the balance of control objectives and safety constraints. Recently, diffusion models have demonstrated a strong ability to model high-dimensional state spaces, giving them an advantage over recent deep learning and reinforcement learning-based methods in complex control tasks. However, they do not inherently address safety concerns. In contrast, while safe reinforcement learning methods consider safety, they typically fail to provide guarantees for satisfying safety constraints. To address these limitations, we propose Safe Conformal Physical system control (SafeConPhy), which optimizes the diffusion model with a provable safety bound iteratively to satisfy the safety constraint. We pre-train a diffusion model on the training set. Given the calibration set and the specific control targets, we derive a provable safety bound using conformal prediction. After iteratively enhancing the safety of the diffusion model with the progressively updated bound, the model’s output can be certified as safe with a user-defined probability. We evaluate our algorithm on two control tasks: 1D Burgers’ equation and 2D incompressible fluid. Our results show that our algorithm satisfies safety constraints, and outperforms prior control methods and safe offline RL algorithms.

1 INTRODUCTION

The control of complex physical systems is critical and essential in many scientific and engineering fields, including fluid dynamics (Hinze & Kunisch, 2001), nuclear fusion (Edwards et al., 1992), and mathematical finance (Soner, 2004). In real-world scenarios, controlling such systems often requires addressing safety concerns (Barros & des Santos, 1998; Argomedo et al., 2013). For example, in fluid dynamics, small errors in control can lead to turbulence or structural damage, while in controlled nuclear fusion, failure to maintain safety constraints could result in catastrophic consequences. Safety, in this context, involves ensuring that the control sequences guide the system to satisfy pre-defined constraints, thereby mitigating risks and preventing hazardous situations (Dawson et al., 2022; Liu et al., 2023a). Notably, safety remains a bottleneck for applying machine learning to specific scientific and engineering problems, as many machine learning algorithms lack the mechanisms to guarantee safety constraints in their control outputs. This gap between performance and safety has become a critical obstacle in deploying machine learning for high-stake applications.

Despite of its importance, the safe control of complex physical systems is challenging. Firstly, to avoid unacceptable risks, one should prevent algorithms without safety guarantees from interacting with the environment, restricting us to an offline setting with pre-collected data. However, the data is often non-optimal and may contain unsafe samples, resulting in a significant gap between the observed data distribution and the near-optimal, safe distribution (Xu et al., 2022; Liu et al., 2023a). Secondly, the algorithm must balance the need for high performance with adherence to safety constraints (Liu et al., 2023a; Zheng et al., 2024).

*Equal contribution. [§]Work done as an intern at Westlake University. [†]Corresponding author.

Table 1: **Comparison between previous deep learning-based control algorithms and our proposed SafeConPhy.** SafeConPhy considers the safety constraints in the control of complex physical systems, and its safety is certifiable before interacting with the environment.

Methods	Complex Physical System	Safety Constraint	Certifiable
DiffPhyCon (Wei et al., 2024)	✓	✗	✗
CDT (Liu et al., 2023b)	✗	✓	✗
TREBI (Lin et al., 2023)	✗	✓	✗
SafeConPhy (ours)	✓	✓	✓

After many years of research on traditional control algorithms (Li et al., 2006; Protas, 2008), the advancement of neural networks leads to the emergence of numerous deep learning-based algorithms (Farahmand et al., 2017; Holl et al., 2020; Hwang et al., 2022). For complex physical systems, which are highly nonlinear and high-dimensional, the deep learning-based methods achieve outstanding results (Hwang et al., 2022; Holl et al., 2020; Wei et al., 2024). However, the above deep learning-based methods generally do not account for safety considerations. Regarding safe offline reinforcement learning (RL), on the one hand, RL methods struggle to optimize long-term control sequences under the constraints of system dynamics (Wei et al., 2024). On the other hand, recent TREBI (Lin et al., 2023) and FISOR (Zheng et al., 2024) utilize diffusion models for planning and theoretically analyzing how to satisfy safety constraints, but they fail to compute the probabilistic bound of safety costs concretely. This limitation prevents their capability of certifying safety before testing, which is inconsistent with the goal of satisfying safety constraints using offline data.

To address these problems, we propose Safe Conformal Physical system control (SafeConPhy), an iterative safety improvement method with a certifiable safety bound. Firstly, in offline settings, the training data are often sub-optimal and unsafe, exhibiting a significant deviation from the desired distribution, which is optimal and safe. Inspired by concepts from conformal prediction (Vovk et al., 2005b; Tibshirani et al., 2019), we estimate the model prediction error under distribution shift based on a portion of split-out training data (called *calibration set*) and the specific control targets. With the estimated prediction error, we compute a probabilistic upper bound of the safety score, and the safety score for the model’s interaction with the environment will be within the upper bound with a user-defined probability. Thus, the model’s safety can be certified by verifying whether the upper bound satisfies the safety constraint. Secondly, we implement a process to improve the model safety by leveraging the upper bound through guidance and fine-tuning iteratively. The guidance step directs the model to stochastically generate multiple samples that potentially satisfy the safety constraint, while the fine-tuning step updates the model by incorporating these samples and the safety bound. This safety improvement process is iterative and continues until the safety upper bound meets the safety constraints.

In summary, the advantages of SafeConPhy are highlighted in Table 1. Our main contributions are as follows: **(1)** We introduce safety constraints into the deep learning-based control of complex physical systems, develop two datasets for safe physical system control tasks to evaluate different methods, and propose the offline algorithm SafeConPhy. **(2)** Considering the model’s prediction error, we provide a certifiable upper bound of the safety score and design an iterative safety improvement process that uses the upper bound to promote the output distribution becoming more optimal and safer. **(3)** We conduct experiments on 1D Burgers’ Equation and 2D incompressible fluid, whose results demonstrate that SafeConPhy can meet the safety constraints and reach better control objectives at the same time.

2 PRELIMINARY

2.1 PROBLEM SETUP

We consider the following safe control problem of complex physical systems:

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \mathcal{J}(\mathbf{u}, \mathbf{w}) \quad \text{s.t.} \quad \mathcal{C}(\mathbf{u}, \mathbf{w}) = 0, \quad s(\mathbf{u}) \leq s_0, \quad (1)$$

where $\mathbf{u}(t, \mathbf{x}) : [0, T] \times \Omega \mapsto \mathbb{R}^{d_{\mathbf{u}}}$ is the system’s state trajectory with dimension $d_{\mathbf{u}}$ and $\mathbf{w}(t, \mathbf{x}) : [0, T] \times \Omega \mapsto \mathbb{R}^{d_{\mathbf{w}}}$ is the external control signal with dimension $d_{\mathbf{w}}$. They are both defined on the

time range $[0, T] \subset \mathbb{R}$ and spatial domain $\Omega \subset \mathbb{R}^D$. $\mathcal{J}(\mathbf{u}, \mathbf{w})$ is the objective of the control problem, and $\mathcal{C}(\mathbf{u}, \mathbf{w}) = 0$ is the physical constraint, such as the partial differential equation. As for the safety constraint, $s(\mathbf{u})$ is the safety score and s_0 is the bound of the safety score. We need to minimize the control objective while satisfying physical constraints and constraining the safety score to stay below the bound, which requires a careful balance between safety and performance. However, it is important to note that safety and performance are not on equal footing, and the pursuit of a better objective should be built upon ensuring safety.

2.2 DIFFUSION MODELS AND DIFFUSION CONTROL

Diffusion models (Ho et al., 2020) learn data distribution from data in a generative way. They present impressive performance in a broad range of generation tasks. Diffusion models involve diffusion/denoising processes: the diffusion process $q(\mathbf{x}^{k+1}|\mathbf{x}^k) = \mathcal{N}(\mathbf{x}^{k+1}; \sqrt{\alpha_k}\mathbf{x}^k, (1 - \alpha_k)\mathbf{I})$ corrupts the data distribution $p(\mathbf{x}_0)$ to a prior distribution $\mathcal{N}(\mathbf{0}, \mathbf{I})$, and the denoising process $p_\theta(\mathbf{x}^{k-1}|\mathbf{x}^k) = \mathcal{N}(\mathbf{x}^{k-1}; \mu_\theta(\mathbf{x}^k, k), \sigma_k^2\mathbf{I})$ makes sampling in a reverse direction. Here k is the diffusion/denoising step, $\{\alpha_k\}_{k=1}^K$ and $\{\sigma_k\}_{k=1}^K$ are the noise and variance schedules. In practice, a denoising network ϵ_θ is trained to estimate the noise to be removed in each step. During inference, the iterative application of ϵ_θ from the prior distribution could generate a new sample that approximately follows the data distribution $p(\mathbf{x})$.

Recently, DiffPhyCon (Wei et al., 2024) applies diffusion models to solve the control problem as in Eq. 1 without the safety constraint $s(\mathbf{u}) \leq s_0$. For brevity, we only summarize its light version. It transforms the physical constraint to a parameterized energy-based model (EBM) $E_\theta(\mathbf{u}, \mathbf{w})$ with the correspondence $p(\mathbf{u}, \mathbf{w}) \propto \exp(-E_\theta(\mathbf{u}, \mathbf{w}))$. Then the problem is converted to an unconstrained optimization over \mathbf{u} and \mathbf{w} for all physical time steps simultaneously:

$$\mathbf{u}^*, \mathbf{w}^* = \arg \min_{\mathbf{u}, \mathbf{w}} [E_\theta(\mathbf{u}, \mathbf{w}) + \lambda \cdot \mathcal{J}(\mathbf{u}, \mathbf{w})], \quad (2)$$

where λ is a hyperparameter. To optimize E_θ , a denoising network ϵ_θ is trained to approximate $\nabla_{\mathbf{u}, \mathbf{w}} E_\theta(\mathbf{u}, \mathbf{w})$ by the following loss:

$$\mathcal{L} = \mathbb{E}_{k \sim U(1, K), (\mathbf{u}, \mathbf{w}) \sim p(\mathbf{u}, \mathbf{w}), \epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})} [\|\epsilon - \epsilon_\theta(\sqrt{\alpha_k}[\mathbf{u}, \mathbf{w}] + \sqrt{1 - \alpha_k}\epsilon, k)\|_2^2], \quad (3)$$

where $\bar{\alpha}_k := \prod_{i=1}^k \alpha_i$. After ϵ_θ is trained, Eq. 2 can be optimized by sampling from an initial sample $(\mathbf{u}^K, \mathbf{w}^K) \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$, and iteratively running the following process

$$(\mathbf{u}^{k-1}, \mathbf{w}^{k-1}) = (\mathbf{u}^k, \mathbf{w}^k) - \eta (\epsilon_\theta([\mathbf{u}^k, \mathbf{w}^k], k) + \lambda \nabla_{\mathbf{u}, \mathbf{w}} \mathcal{G}(\hat{\mathbf{u}}^k, \hat{\mathbf{w}}^k)) + \xi, \quad \xi \sim \mathcal{N}(\mathbf{0}, \sigma_k^2 \mathbf{I}) \quad (4)$$

under the guidance of $\mathcal{G} = \mathcal{J}$ for $k = K, K-1, \dots, 1$. Here $[\hat{\mathbf{u}}^k, \hat{\mathbf{w}}^k]$ is the noise-free estimation of $[\mathbf{u}^0, \mathbf{w}^0]$. The final sampling step yields the solution \mathbf{w}^0 for the optimization problem in Eq. 2.

3 METHOD

In this section, we introduce our proposed method SafeConPhy, with its overall framework outlined in Figure 1. First, in Section 3.1, we briefly outline the overall steps of the algorithm. Next, in Section 3.2, we explain how conformal prediction is applied to estimate the safety score s under distribution shift in a probabilistic sense, and theoretically derive the formula for the provable safety bound s_+ . Finally, in Section 3.3, we detail the implementation of the entire algorithm. Specifically, we first describe the detailed implementation of this estimation. Additionally, we describe how the estimated safety score is utilized to modify the distribution of generated control sequences through guidance and fine-tuning.

3.1 OVERALL PROCEDURES

In this section, we present the workflow of our algorithm as in Figure 1. We set aside a portion of the original training data as the *calibration set* D_{cal} , which will be used later to estimate the model’s prediction errors. The remaining data, which will be used for actual training, is referred to as the training data D_{train} . After pre-training with D_{train} as described in Eq. 3, we get the diffusion model p_θ which models the joint distribution of $[\mathbf{u}, \mathbf{w}]$.




Figure 1: **Overview of SafeConPhy.** First, we pre-train a diffusion model p_θ on the training data. Then, we derive a safety bound to certify that the model satisfies the safety constraints. To satisfy the safety constraint, we further design a loss function term based on the safety bound to optimize the diffusion model.

Algorithm 1 Inference of SafeConPhy

- 1: **Require** Calibration set D_{cal} , training set D_{train} , confidence level α , control objective $\mathcal{J}(\cdot)$, safety score $s(\cdot)$, number of iterations N
 - 2: **for** $n = 1, \dots, N$ **do**
 - 3: Compute the weighted score set $\tilde{\mathcal{S}}$ with D_{cal} // Eq. 9
 - 4: Get the quantile $Q(1 - \alpha; \tilde{\mathcal{S}})$
 - 5: Sample the control sequence \mathbf{w} with guidance \mathcal{G} // Eq. 11
 - 6: Compute $s_+(\tilde{\mathbf{u}}_\theta(\mathbf{w}))$ with conditionally sampled $\tilde{\mathbf{u}}_\theta(\mathbf{w})$ // Eq. 10
 - 7: Take gradient descent step on $\nabla_\theta \mathcal{L}_{\text{fine-tune}}$ // Eq. 12
 - 8: **end for**
 - 9: Sample the control sequence \mathbf{w} with guidance \mathcal{G} // Eq. 11
 - 10: **return** \mathbf{w}
-

Next, we conduct Iterative Safety Improvement, where we apply the provable safety bound s_+ under distribution shift to enhance model safety iteratively. First, the calculation of the safety bound s_+ runs throughout the entire loop. The safety bound basically takes the calibration set to obtain a corresponding set of model prediction errors (called the *score set*). Furthermore, we take into account the distribution shift between the calibration set and data generated based on control targets, so we apply weighting to the set of model prediction errors to get the *weighted score set* as Eq. 9. Then, we use the quantile of this set to represent the error in the model’s predicted safety score.

Second, ‘iterative’ refers to the process where we cyclically use the guidance to generate samples, and then combine these samples with model prediction errors to fine-tune the model parameters, thereby improving the model’s safety. Specifically, in each iteration, we sample the control sequence \mathbf{w} under the guidance as described in Eq. 11 containing s_+ . After getting \mathbf{w} , we conditionally sample $\tilde{\mathbf{u}}_\theta(\mathbf{w})$ to get the provable safety bound $s_+(\tilde{\mathbf{u}}_\theta(\mathbf{w}))$. And we can now take the fine-tuning loss $\mathcal{L}_{\text{fine-tune}}$ involving the training data D_{train} and also the progressively updated $s_+(\tilde{\mathbf{u}}_\theta(\mathbf{w}))$ as in Eq. 12 to fine-tune the model parameters θ .

Finally, after several iterations, we use the fine-tuned diffusion model, once again under the influence of guidance, to generate the control sequences, which serve as the final output of the algorithm. The complete inference process can be seen in Algorithm 1.

3.2 PROVABLE SAFETY BOUND WITH CONFORMAL PREDICTION

In offline safety control problems, the gap between pre-collected data and the target distribution exacerbates models’ prediction errors, which can be critical in ensuring safety. To address this issue, we employ the conformal prediction technique to obtain a provable safety bound, ensuring that the true safety score is included within this estimate with a provable level of confidence, without requiring additional assumptions about the model or the data distribution.

Calibration set and score set. To achieve the goal mentioned above, we first set aside a portion of the training dataset, which is not used for training, as the *calibration set* D_{cal} . After training, we

take the calibration set to obtain the *score set*, which is defined as

$$\mathcal{S} := \{|s(\tilde{\mathbf{u}}_\theta(\mathbf{w}_i)) - s(\mathbf{u}_i)| : (\mathbf{u}_i, \mathbf{w}_i) \in D_{\text{cal}}\}. \quad (5)$$

Here \mathbf{u}_i and \mathbf{w}_i represent different samples, $\tilde{\mathbf{u}}_\theta(\mathbf{w})$ is the system state conditioned on control \mathbf{w}_i and predicted by the model, and θ is the parameters of the model. The score set can be considered as recording the estimation errors of the model with respect to the safety score s .

Weighted score set under distribution shift. However, since the data distribution in the calibration set differs from the final model-generated data distribution used for control, it does not satisfy the exchangeability¹ condition required by conformal prediction (Vovk et al., 2005a; Papadopoulos et al., 2002; Lei et al., 2016). Intuitively, each sample in the calibration set has a different probability of appearing in the final data distribution generated by the model, so we further apply weighting to the elements of the score set.

We define $p(\mathbf{u}, \mathbf{w})$ as the distribution of the calibration set. And we let $\tilde{p}(\mathbf{u}, \mathbf{w})$ be the distribution of the test data, where \mathbf{w} is generated by the model based on the control task and safety constraints, and \mathbf{u} is the true system state obtained from the interaction between the control sequence \mathbf{w} and the environment. According to conformal prediction under covariate shift (Tibshirani et al., 2019), the calculation of the weights is as follows:

$$\omega(\mathbf{u}_i, \mathbf{w}_i) := \frac{d\tilde{p}(\mathbf{u}_i, \mathbf{w}_i)}{dp(\mathbf{u}_i, \mathbf{w}_i)} = \frac{d\tilde{p}(\mathbf{w}_i)d\tilde{p}(\mathbf{u}_i|\mathbf{w}_i)}{dp(\mathbf{w}_i)dp(\mathbf{u}_i|\mathbf{w}_i)}, \quad (6)$$

where dp denotes the probability density function of distribution p . Since $d\tilde{p}(\mathbf{u}_i|\mathbf{w}_i)$ and $dp(\mathbf{u}_i|\mathbf{w}_i)$ both represent the physical constraints of the system itself, the weights can be simplified as $\omega(\mathbf{u}_i, \mathbf{w}_i) = \frac{d\tilde{p}(\mathbf{w}_i)}{dp(\mathbf{w}_i)}$. We note that as described in Eq. 11, \mathbf{w}_i is generated by the energy-based model under guidance \mathcal{G} . In the safe control problems, guidance \mathcal{G} encompasses both the control objective and safety constraints and will be detailed in Section 3.3. With the influence of \mathcal{G} , $\tilde{p}(\mathbf{u}_i, \mathbf{w}_i) \propto \exp(-E_\theta(\mathbf{u}_i, \mathbf{w}_i) - \mathcal{G}(\mathbf{u}_i, \mathbf{w}_i)) \propto p_\theta(\mathbf{u}_i, \mathbf{w}_i) \exp(-\mathcal{G}(\mathbf{u}_i, \mathbf{w}_i))$, where p_θ is distribution learned by the diffusion model. Thus we obtain

$$\omega(\mathbf{u}_i, \mathbf{w}_i) = C \frac{dp_\theta(\mathbf{u}_i, \mathbf{w}_i)e^{-\mathcal{G}(\mathbf{u}_i, \mathbf{w}_i)}}{dp(\mathbf{u}_i, \mathbf{w}_i)}, \quad (7)$$

where C is a constant. Given that the calibration set and the training dataset follow the same distribution, and assuming that the impact of the diffusion model's learning error on the dataset is sufficiently small relative to the second term $e^{-\mathcal{G}(\mathbf{u}_i, \mathbf{w}_i)}$, we can approximate the weight as $\omega(\mathbf{u}_i, \mathbf{w}_i) = Ce^{-\mathcal{G}(\mathbf{u}_i, \mathbf{w}_i)}$. Finally, we normalize the weight and obtain

$$\hat{\omega}(\mathbf{u}_i, \mathbf{w}_i) = \frac{Ce^{-\mathcal{G}(\mathbf{u}_i, \mathbf{w}_i)}}{\sum_{(\mathbf{u}_j, \mathbf{w}_j) \in D_{\text{cal}}} Ce^{-\mathcal{G}(\mathbf{u}_j, \mathbf{w}_j)}} = \frac{e^{-\mathcal{G}(\mathbf{u}_i, \mathbf{w}_i)}}{\sum_{(\mathbf{u}_j, \mathbf{w}_j) \in D_{\text{cal}}} e^{-\mathcal{G}(\mathbf{u}_j, \mathbf{w}_j)}}. \quad (8)$$

The *weighted score set* is then defined as

$$\tilde{\mathcal{S}} := \{\hat{\omega}(\mathbf{u}_i, \mathbf{w}_i)|s(\tilde{\mathbf{u}}_\theta(\mathbf{w}_i)) - s(\mathbf{w}_i)| : (\mathbf{u}_i, \mathbf{w}_i) \in D_{\text{cal}}\}. \quad (9)$$

Upper bound s_+ on the confidence level α . For a given control sequence \mathbf{w} , we provide an upper bound s_+ below, such that with at least $1 - \alpha$ probability, the true s is smaller than s_+ . In detail, we exploit the weighted score set to define the s_+ as

$$s_+(\tilde{\mathbf{u}}_\theta(\mathbf{w})) = s(\tilde{\mathbf{u}}_\theta(\mathbf{w})) + Q(1 - \alpha; \tilde{\mathcal{S}}), \quad (10)$$

where $Q(1 - \alpha; \tilde{\mathcal{S}})$ is $\text{Quantile}((1 - \alpha)(1 + \frac{1}{|D_{\text{cal}}|}); \tilde{\mathcal{S}})$ and is still a differentiable function with respect to θ , and $|D_{\text{cal}}|$ is the cardinality of D_{cal} . The precise and formal meaning of the probabilistic upper bound is demonstrated through the following lemma.

Lemma 1. Assume samples $(\mathbf{u}_i, \mathbf{w}_i) \sim p$ in the calibration set are independent, and the test set $(\mathbf{u}, \mathbf{w}) \sim \tilde{p}$ is also independent with the calibration set. Assume p is absolutely continuous with respect to \tilde{p} , s_+ is defined as in Eq. 10, then $\mathbb{P}[s(\mathbf{u}(\mathbf{w})) \leq s_+(\tilde{\mathbf{u}}_\theta(\mathbf{w}))] \geq 1 - \alpha$, where $\mathbf{u}(\mathbf{w})$ is the real system state corresponding to the control sequence \mathbf{w} , and $\tilde{\mathbf{u}}_\theta(\mathbf{w})$ is the system state predicted by the model conditioned on the same control.

¹ $(\mathbf{w}_i, s_i)_{i=1}^N$ are exchangeable if, for any permutation σ of $[1, N]$, $\mathcal{P}((\mathbf{u}_1, s_1), \dots, (\mathbf{u}_N, s_N)) = \mathcal{P}((\mathbf{u}_{\sigma(1)}, s_{\sigma(1)}), \dots, (\mathbf{u}_{\sigma(N)}, s_{\sigma(N)}))$, where \mathcal{P} is the joint distribution.

3.3 TARGET CONTROL GENERATION BASED ON ESTIMATED SAFETY

Next, based on the deduced s_+ , we describe the detailed implementation of modules in SafeConPhy.

Conditionally sample $\tilde{\mathbf{u}}_\theta(\mathbf{w})$. In our proposed algorithm, we need to sample from the conditional distribution $p(\mathbf{u}|\mathbf{w})$ with the model that learns the joint distribution $p(\mathbf{u}, \mathbf{w})$. To achieve this, at each denoising step of the sampling process, we replace the noisy \mathbf{w} in the input of the denoising network with the actual clean \mathbf{w} that serves as the condition (Chung et al., 2023). In fact, this situation represents a special case of the data distribution that the denoising network encounters during training, where the \mathbf{u} part is noisy, while the \mathbf{w} part remains noise-free.

Guidance \mathcal{G} . Guidance is the first method we adopt to steer the model’s output toward satisfying both the control objectives and safety constraints. It plays a role during the sampling process of the diffusion model. When considering safety, the specific form of our guidance is as follows:

$$\mathcal{G}(\mathbf{u}, \mathbf{w}) = \mathcal{J}(\mathbf{u}, \mathbf{w}) + \gamma \max[s_+(\mathbf{u}(\mathbf{w})) - s_0, 0]. \quad (11)$$

Note that although we follow the previous symbol $s_+(\mathbf{u}(\mathbf{w}))$, during sampling, \mathbf{u} and \mathbf{w} are actually generated by the diffusion model jointly but not conditionally. The specific denoising step of implementing guidance follows Eq. 4.

Fine-tuning. The second method for adjusting the output data distribution of the model is fine-tuning, which achieves the adjustment by optimizing the model parameters θ . Specifically, both terms in s_+ , as shown in Eq. 10, are functions of θ . Therefore, it is both reasonable and effective to compute the gradient of s_+ with respect to θ to take the gradient descent step. It is worth noting that retaining the computation graph for all denoising steps with respect to θ would result in an unmanageable memory overhead. Therefore, when we need to keep the computation graph during denoising for gradient calculation, we only retain the computation graph of the final denoising step.

To optimize the safety score and the diffusion loss (referred to Eq. 3) simultaneously, we form the fine-tune loss $\mathcal{L}_{\text{fine-tune}}$ as the weighted sum of both the safety loss $\mathcal{L}_{\text{safe}}$ and diffusion loss $\mathcal{L}_{\text{diffusion}}$:

$$\begin{aligned} \mathcal{L}_{\text{fine-tune}} &= \mathcal{L}_{\text{safe}} + \beta \mathcal{L}_{\text{diffusion}} \\ &= \sum_{\mathbf{w} \in D_{\text{sampled}}} \max[s_+(\tilde{\mathbf{u}}_\theta(\mathbf{w})) - s_0, 0] + \beta \sum_{(\mathbf{u}, \mathbf{w}) \in D_{\text{train}}} \|\epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_k}[\mathbf{u}, \mathbf{w}] + \sqrt{1 - \bar{\alpha}_k}\epsilon, \mathbf{u}_0, k)\|_2^2, \end{aligned} \quad (12)$$

where \mathbf{w} in the first term is from D_{sampled} sampled according to the guidance described above, (\mathbf{u}, \mathbf{w}) in the second term is from the training set D_{train} , k is the denoising step, $\epsilon \sim \mathcal{N}(0, \mathbf{I})$, $\bar{\alpha}_k := \prod_{i=1}^k \alpha_i$ is the product of noise schedules and $[\mathbf{u}, \mathbf{w}]$ means the concatenation of \mathbf{u} and \mathbf{w} .

4 EXPERIMENT

To verify our statements that SafeConPhy can both achieve safety and reach lower control objectives than other methods, we conduct experiments on safe offline control problems on 1D Burgers’ equation and 2D incompressible fluid. Besides, to evaluate the quality of safety, for different problems, we provide several corresponding metrics.

For comparison, we choose Behavior Cloning (BC) (Pomerleau, 1988), and safe reinforcement and imitation learning methods involving BC with safe data filtering (BC-Safe), Constrained Decision Transformer (CDT) (Liu et al., 2023b) and CDT with safe data filtering (CDT-Safe), diffusion-based method TREBI (Lin et al., 2023). Note that CDT shows the best performance in the Offline Safe RL benchmark OSRL (Liu et al., 2023a). In addition, we combine the physical system control method Supervised Learning (Hwang et al., 2022) with the Lagrangian approach (Chow et al., 2018) (SL-Lag) to enforce safety constraints. We also apply the classical control method PID (Li et al., 2006). We provide the anonymous code here.

4.1 1D BURGERS’ EQUATION

Experiment settings. 1D Burgers’ equation is a fundamental equation that governs various physical systems including fluid dynamics and gas dynamics. Here we follow previous works (Hwang et al., 2022; Mowlavi & Nabi, 2023) and consider the Dirichlet boundary condition along with an external

Table 2: **Results of 1D Burgers’ equation.** Gray: s_{norm} is greater than 1 (unsafe). Black: s_{norm} is smaller than 1 (safe). **Bold:** Safe trajectories with *lowest* \mathcal{J} .

Methods	$\mathcal{J} \downarrow$	$s_{\text{norm}} \downarrow$	$\mathcal{R}_{\text{sample}} \downarrow$	$\mathcal{R}_{\text{time}} \downarrow$	$\mathcal{R}_{\text{point}} \downarrow$
BC	0.0001	1.9954	38%	13%	1.2%
BC-Safe	0.0002	1.9601	14%	3%	0.2%
PID	0.0968	0.5691	0%	0%	0.0%
SL-Lag	0.0115	0.6817	0%	0%	0.0%
CDT	0.0026	1.9220	8%	1%	0.1%
CDT-Safe	0.0021	0.8570	0%	0%	0.0%
TREBI	0.0074	0.7821	0%	0%	0.0%
SafeConPhy (ours)	0.0011	0.8388	0%	0%	0.0%

force $\mathbf{w}(t, x)$. This equation is formulated as follows:

$$\begin{cases} \frac{\partial \mathbf{u}(t, x)}{\partial t} = -\mathbf{u}(t, x) \cdot \frac{\partial \mathbf{u}(t, x)}{\partial x} + \nu \frac{\partial^2 \mathbf{u}(t, x)}{\partial x^2} + \mathbf{w}(t, x) & \text{in } [0, T] \times \Omega \\ \mathbf{u}(t, x) = 0 & \text{on } [0, T] \times \partial\Omega \\ \mathbf{u}(0, x) = \mathbf{u}_0(x) & \text{in } \{t = 0\} \times \Omega, \end{cases} \quad (13)$$

where ν denotes the viscosity parameter, while \mathbf{u}_0 signifies the initial condition. We set $\nu = 0.01$, $T = 1$ and $\Omega = [0, 1]$. Given a target state $\mathbf{u}_d(x)$, the primary control objective \mathcal{J} is to minimize the control error between the final state \mathbf{u}_T and the target state \mathbf{u}_d .

$$\mathcal{J} := \int_{\Omega} |\mathbf{u}(T, x) - \mathbf{u}_d(x)|^2 dx. \quad (14)$$

We define the safety score as: $s(\mathbf{u}) := \sup_{(t, x) \in [0, T] \times \Omega} \{\mathbf{u}(t, x)^2\}$. Considering the safety constraint s_0 , if $s(\mathbf{u}) > s_0$, the state trajectory \mathbf{u} is unsafe, and if $s(\mathbf{u}) \leq s_0$, the state trajectory \mathbf{u} is safe. The bound of safety score s_0 is set to 0.64 in our experiment. According to this bound, 89.7% of samples are unsafe among the training set, 90% of samples are unsafe among the calibration set and all of the samples in the test set are unsafe. More details can be found in Appendix F.1.

To better evaluate whether the set of state trajectories controlled by the model is safe, we define the normalized safety score:

$$s_{\text{norm}} := \frac{1}{|\mathcal{N}_1|} \sum_{i \in \mathcal{N}_1} \frac{s(\mathbf{u}_i)}{s_0} + \frac{1}{|\mathcal{N}_2|} \sum_{i \in \mathcal{N}_2} \frac{s(\mathbf{u}_i)}{s_0}, \quad (15)$$

where $\mathcal{N}_1 = \{i \mid \mathbf{u}_i \leq s_0\}$, $\mathcal{N}_2 = \{i \mid \mathbf{u}_i > s_0\}$. Note that $s(\mathbf{u})$ and s_0 are non-negative. If the state trajectories are all safe, s_{norm} is smaller than 1; If *any* state trajectory is unsafe, s_{norm} is greater than 1. Additionally, we compute three unsafe rates to assess the safety levels of different methods’ control results. $\mathcal{R}_{\text{sample}}$ denotes the proportion of unsafe trajectories among total trajectories²; $\mathcal{R}_{\text{time}}$ denotes the proportion of unsafe timesteps among all timesteps; $\mathcal{R}_{\text{point}}$ denotes the proportion of unsafe spatial lattice points in all spatial lattice points across all time steps.

Results. In Table 2, We report the results of the control objective \mathcal{J} , the safety score s_{norm} and other safe metrics of different methods. SafeConPhy can meet the safety constraint and achieve the best control objective at the same time. As shown in Figure 2, given the initial condition and the final state (control target), SafeConPhy can control a state trajectory that satisfies the safety constraint and control target. Other methods either suffer from constraint violations or suboptimal objectives.

4.2 2D INCOMPRESSIBLE FLUID

Experiment settings. We then consider the control problems of 2D incompressible fluid, which follows the 2D Navier-Stokes equation. The control task we consider is to maximize the amount of smoke that passes through the target bucket in the fluid flow with obstacles and openings, while constraining the amount of smoke passing through the dangerous region under the safety bound (Wei et al., 2024). Specifically, referring to Figure 3, the control objective \mathcal{J} is defined as the negative rate of smoke passing through the target bucket located at the center top, while the safety score

²If any point in the full trajectory is unsafe, this trajectory is unsafe. So $\mathcal{R}_{\text{sample}}$ is the most stringent metric.

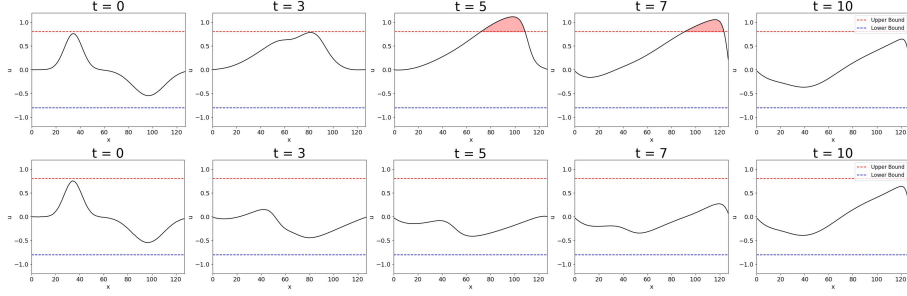


Figure 2: **Visualizations of the 1D Burgers' equation.** The top row shows the original trajectory corresponding to the control target, and the bottom row is the trajectory controlled by SafeConPhy.




Figure 3: **Visualization of the 2D incompressible fluid control problem.**

s corresponds to the rate of smoke entering the hazardous red region. It is important to note that there is a trade-off between controlling the flow through the hazardous region and achieving a more optimal control objective, which imposes higher demands on the algorithm. We set the safety score bound to $s_0 = 0.1$.

Moreover, this control task is particularly challenging: not only does it require indirect control, which means that control can only be applied to the peripheral region, but the spatial control parameters reach as many as 1,792. Among all the training data, 53.1% of the samples are unsafe. The average safety score of the dataset is 0.3215. Other details can be found in Appendix G.1.

Results. We report results in Table 3. Here PID is inapplicable and SL-Lag fails to achieve reasonable control results. Due to the challenges of the task, no method can guarantee that all samples meet the safety requirements, so we introduce additional metrics to assess safety. Here we define the normalized safety score s_{norm} as s/s_0 and define \mathcal{R} as the rate of unsafe samples. Additionally, we introduce another metric $\max[s - s_0, 0]$. When s does not exceed the bound s_0 , this metric is 0. If s exceeds s_0 , the metric reflects the amount by which it is surpassed. From the results, we can see that our method successfully keeps s_{norm} below 1, and other safety metrics are also comparable to other safe baselines ($s_{\text{norm}} \leq 1$). Besides, among safe methods, SafeConPhy achieves the lowest \mathcal{J} , even reaching control performance similar to methods that do not consider safety like BC.

5 CONCLUSION

In this paper, we introduce Safe Conformal Physical system control (SafeConPhy), a probabilistic generative method for safe control physical systems problems. Targeting the meaningful and important offline setting, we provide a provable probabilistic upper bound of the safety score. We then perform guidance and fine-tuning with this provable safety bound iteratively, improving safety and certifying it with a user-defined probability. Experiment results on 1D Burgers' equation and 2D incompressible fluid demonstrate that on the basis of satisfying safety constraints, SafeConPhy is able to achieve better control objectives. We believe that our method is beneficial for making deep learning-based physical control safer, improving the trustworthiness for deploying to the real world.

Table 3: **2D incompressible fluid control results.** Gray: s_{norm} is greater than 1 (unsafe). Black: s_{norm} is smaller than 1 (safe). **Bold:** methods marked in black with lowest \mathcal{J} .

Methods	$\mathcal{J} \downarrow$	$s_{\text{norm}} \downarrow$	$\max[s - s_0, 0] \downarrow$	$\mathcal{R} \downarrow$
BC	-0.7125	7.3402	0.7160	88%
BC-Safe	-0.2520	0.3463	0.0330	8%
CDT	-0.7133	3.0778	0.2726	34%
CDT-Safe	-0.6360	0.5073	0.0292	18%
TREBI	-0.6105	0.9096	0.0537	30%
SafeConPhy (ours)	-0.7035	0.6092	0.0380	14%

REFERENCES

- Joshua Achiam, David Held, Aviv Tamar, and Pieter Abbeel. Constrained policy optimization. In *International conference on machine learning*, pp. 22–31. PMLR, 2017.
- Federico Bribiesca Argomedo, Emmanuel Witrant, Christophe Prieur, Sylvain Bremond, Remy Nouailletas, and Jean-Francois Artaud. Lyapunov-based distributed control of the safety-factor profile in a tokamak plasma. *Nuclear Fusion*, 53(3):033005, 2013.
- E. Barros and M.V.D. des Santos. A safe, accurate intravenous infusion control system. *IEEE Micro*, 18(5):12–21, 1998. doi: 10.1109/40.735940.
- Lili Chen, Kevin Lu, Aravind Rajeswaran, Kimin Lee, Aditya Grover, Misha Laskin, Pieter Abbeel, Aravind Srinivas, and Igor Mordatch. Decision transformer: Reinforcement learning via sequence modeling. *Advances in neural information processing systems*, 34:15084–15097, 2021.
- Victor Chernozhukov, Kaspar Wüthrich, and Zhu Yinchu. Exact and robust conformal inference methods for predictive machine learning with dependent data. In *Conference On learning theory*, pp. 732–749. PMLR, 2018.
- Yinlam Chow, Mohammad Ghavamzadeh, Lucas Janson, and Marco Pavone. Risk-constrained reinforcement learning with percentile risk criteria. *Journal of Machine Learning Research*, 18(167):1–51, 2018.
- Hyungjin Chung, Jeongsol Kim, Michael Thompson Mccann, Marc Louis Klasky, and Jong Chul Ye. Diffusion posterior sampling for general noisy inverse problems. In *The Eleventh International Conference on Learning Representations*, 2023. URL <https://openreview.net/forum?id=OnD9zgAGT0k>.
- Charles Dawson, Zengyi Qin, Sicun Gao, and Chuchu Fan. Safe nonlinear control using robust neural lyapunov-barrier functions. In *Conference on Robot Learning*, pp. 1724–1735. PMLR, 2022.
- Jonas Degraeve, Federico Felici, Jonas Buchli, Michael Neunert, Brendan Tracey, Francesco Carpanese, Timo Ewalds, Roland Hafner, Abbas Abdolmaleki, Diego de Las Casas, et al. Magnetic control of tokamak plasmas through deep reinforcement learning. *Nature*, 602(7897):414–419, 2022.
- Prafulla Dhariwal and Alex Nichol. Diffusion Models Beat GANs on Image Synthesis, June 2021. URL <http://arxiv.org/abs/2105.05233>. arXiv:2105.05233 [cs, stat].
- Robert M Edwards, Kwang Y Lee, and Asok Ray. Robust optimal control of nuclear reactors and power plants. *Nuclear Technology*, 98(2):137–148, 1992.
- Amir-massoud Farahmand, Saleh Nabi, and Daniel N. Nikovski. Deep reinforcement learning for partial differential equation control. In *2017 American Control Conference (ACC)*, pp. 3120–3127, 2017. doi: 10.23919/ACC.2017.7963427.
- Haodong Feng, Yue Wang, Hui Xiang, Zhiyang Jin, and Dixia Fan. How to control hydrodynamic force on fluidic pinball via deep reinforcement learning. *Physics of Fluids*, 35(4), 2023.
- Haodong Feng, Dehan Yuan, Jiale Miao, Jie You, Yue Wang, Yi Zhu, and Dixia Fan. Efficient navigation of a robotic fish swimming across the vortical flow field. *arXiv preprint arXiv:2405.14251*, 2024.
- Isaac Gibbs and Emmanuel Candes. Adaptive conformal inference under distribution shift. *Advances in Neural Information Processing Systems*, 34:1660–1672, 2021.
- Dan Hendrycks, Mantas Mazeika, Duncan Wilson, and Kevin Gimpel. Using trusted data to train deep networks on labels corrupted by severe noise. *Advances in neural information processing systems*, 31, 2018.
- Michael Hinze and Karl Kunisch. Second order methods for optimal control of time-dependent fluid flow. *SIAM Journal on Control and Optimization*, 40(3):925–946, 2001.

- Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. *Advances in neural information processing systems*, 33:6840–6851, 2020.
- Jonathan Ho, Tim Salimans, Alexey Gritsenko, William Chan, Mohammad Norouzi, and David J Fleet. Video diffusion models. *arXiv:2204.03458*, 2022.
- Philipp Holl, Nils Thuerey, and Vladlen Koltun. Learning to control pdes with differentiable physics. In *International Conference on Learning Representations*, 2020.
- Rakhoon Hwang, Jae Yong Lee, Jin Young Shin, and Hyung Ju Hwang. Solving pde-constrained control problems using operator learning. In *Proceedings of the AAAI Conference on Artificial Intelligence*, pp. 4504–4512, 2022.
- Michael Janner, Yilun Du, Joshua Tenenbaum, and Sergey Levine. Planning with diffusion for flexible behavior synthesis. *Proceedings of Machine Learning Research*, 162:9902–9915, 17–23 Jul 2022.
- Michael A Johnson and Mohammad H Moradi. *PID control*. Springer, 2005.
- Jongmin Lee, Cosmin Paduraru, Daniel J Mankowitz, Nicolas Heess, Doina Precup, Kee-Eung Kim, and Arthur Guez. COptDICE: Offline constrained reinforcement learning via stationary distribution correction estimation. In *International Conference on Learning Representations*, 2022. URL <https://openreview.net/forum?id=FLA55mBee6Q>.
- Jing Lei, Max Grazier G’Sell, Alessandro Rinaldo, Ryan J. Tibshirani, and Larry A. Wasserman. Distribution-free predictive inference for regression. *Journal of the American Statistical Association*, 113:1094 – 1111, 2016. URL <https://api.semanticscholar.org/CorpusID:137414119>.
- Yun Li, Kiam Heong Ang, and G.C.Y. Chong. Pid control system analysis and design. *IEEE Control Systems Magazine*, 26(1):32–41, 2006. doi: 10.1109/MCS.2006.1580152.
- Qian Lin, Bo Tang, Zifan Wu, Chao Yu, Shangqin Mao, Qianlong Xie, Xingxing Wang, and Dong Wang. Safe offline reinforcement learning with real-time budget constraints. In *International Conference on Machine Learning*, pp. 21127–21152. PMLR, 2023.
- Zuxin Liu, Zhepeng Cen, Vladislav Isenbaev, Wei Liu, Steven Wu, Bo Li, and Ding Zhao. Constrained variational policy optimization for safe reinforcement learning. In *International Conference on Machine Learning*, pp. 13644–13668. PMLR, 2022.
- Zuxin Liu, Zijian Guo, Haohong Lin, Yihang Yao, Jiacheng Zhu, Zhepeng Cen, Hanjiang Hu, Wenhao Yu, Tingnan Zhang, Jie Tan, et al. Datasets and benchmarks for offline safe reinforcement learning. *arXiv preprint arXiv:2306.09303*, 2023a.
- Zuxin Liu, Zijian Guo, Yihang Yao, Zhepeng Cen, Wenhao Yu, Tingnan Zhang, and Ding Zhao. Constrained decision transformer for offline safe reinforcement learning. In *International Conference on Machine Learning*, pp. 21611–21630. PMLR, 2023b.
- Alnur Ali Maxime Cauchois, Suyash Gupta and John C. Duchi. Robust validation: Confident predictions even when distributions shift. *Journal of the American Statistical Association*, 0(0): 1–66, 2024. doi: 10.1080/01621459.2023.2298037. URL <https://doi.org/10.1080/01621459.2023.2298037>.
- Saviz Mowlavi and Saleh Nabi. Optimal control of pdes using physics-informed neural networks. *Journal of Computational Physics*, 473:111731, 2023.
- Guido Novati, Siddhartha Verma, Dmitry Alexeev, Diego Rossinelli, Wim M Van Rees, and Petros Koumoutsakos. Synchronisation through learning for two self-propelled swimmers. *Bioinspiration & biomimetics*, 12(3):036001, 2017.
- Yangchen Pan, Amir-massoud Farahmand, Martha White, Saleh Nabi, Piyush Grover, and Daniel Nikovski. Reinforcement learning with function-valued action spaces for partial differential equation control. In Jennifer Dy and Andreas Krause (eds.), *Proceedings of the 35th International Conference on Machine Learning*, volume 80, pp. 3986–3995. PMLR, 10–15 Jul 2018.

- Harris Papadopoulos, Kostas Proedrou, Vladimir Vovk, and Alexander Gammerman. Inductive confidence machines for regression. In *European Conference on Machine Learning*, 2002. URL <https://api.semanticscholar.org/CorpusID:42084298>.
- Aleksandr Podkopaev and Aaditya Ramdas. Distribution-free uncertainty quantification for classification under label shift. In Cassio de Campos and Marloes H. Maathuis (eds.), *Proceedings of the Thirty-Seventh Conference on Uncertainty in Artificial Intelligence*, volume 161 of *Proceedings of Machine Learning Research*, pp. 844–853. PMLR, 27–30 Jul 2021. URL <https://proceedings.mlr.press/v161/podkopaev21a.html>.
- Dean A Pomerleau. Alvin: An autonomous land vehicle in a neural network. *Advances in neural information processing systems*, 1, 1988.
- Bartosz Protas. Adjoint-based optimization of pde systems with alternative gradients. *Journal of Computational Physics*, 227(13):6490–6510, 2008.
- Jean Rabault, Miroslav Kuchta, Atle Jensen, Ulysse Réglade, and Nicolas Cerardi. Artificial neural networks trained through deep reinforcement learning discover control strategies for active flow control. *Journal of fluid mechanics*, 865:281–302, 2019.
- H Mete Soner. *Stochastic optimal control in finance*. Scuola normale superiore, 2004.
- Adam Stooke, Joshua Achiam, and Pieter Abbeel. Responsive safety in reinforcement learning by pid lagrangian methods. In *International Conference on Machine Learning*, pp. 9133–9143. PMLR, 2020.
- Ryan J. Tibshirani, Rina Foygel Barber, Emmanuel J. Candès, and Aaditya Ramdas. Conformal prediction under covariate shift. In *Neural Information Processing Systems*, 2019. URL <https://api.semanticscholar.org/CorpusID:115140768>.
- Vladimir Vovk, Alexander Gammerman, and Glenn Shafer. *Algorithmic learning in a random world*. Springer US, United States, 2005a. ISBN 0387001522. doi: 10.1007/b106715.
- Vladimir Vovk, Alexander Gammerman, and Glenn Shafer. *Algorithmic learning in a random world*, volume 29. Springer, 2005b.
- Long Wei, Peiyan Hu, Ruiqi Feng, Haodong Feng, Yixuan Du, Tao Zhang, Rui Wang, Yue Wang, Zhi-Ming Ma, and Tailin Wu. A generative approach to control complex physical systems. *arXiv preprint arXiv:2407.06494*, 2024.
- Haoran Xu, Xianyu Zhan, and Xiangyu Zhu. Constraints penalized q-learning for safe offline reinforcement learning. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 36, pp. 8753–8760, 2022.
- Shuang Zhang, Xinyu Qian, Zhijie Liu, Qing Li, and Guang Li. Pde modeling and tracking control for the flexible tail of an autonomous robotic fish. *IEEE Transactions on Systems, Man, and Cybernetics: Systems*, 52(12):7618–7627, 2022.
- Yiming Zhang, Quan Vuong, and Keith Ross. First order constrained optimization in policy space. *Advances in Neural Information Processing Systems*, 33:15338–15349, 2020.
- Yinan Zheng, Jianxiong Li, Dongjie Yu, Yujie Yang, Shengbo Eben Li, Xianyu Zhan, and Jingjing Liu. Safe offline reinforcement learning with feasibility-guided diffusion model. *arXiv preprint arXiv:2401.10700*, 2024.

A RELATED WORK

A.1 CONTROL OF PHYSICAL SYSTEMS

The development of control methods in physical systems is critical across various engineering areas, including PID (Li et al., 2006), supervised learning (SL) (Holl et al., 2020; Hwang et al., 2022), reinforcement learning (RL) (Farahmand et al., 2017; Pan et al., 2018; Rabault et al., 2019), and physics-informed neural networks (PINNs) (Mowlavi & Nabi, 2023). Among these, PID is one of the earliest and most widely used method (Johnson & Moradi, 2005), known for its simplicity and effectiveness in regulating physical systems; however, it faces challenges in parameter tuning and struggles with highly nonlinear or time-varying systems. With the advancement of deep learning, SL (Holl et al., 2020) has been applied to optimize control sequences through backpropagation over entire trajectories, but it lacks the adaptability to dynamic environments since it is typically trained on fixed datasets. To overcome the above issue, RL enhances adaptability by leveraging diverse datasets or interactions with the environment, achieving notable success in controlling physical systems, including fluid dynamics (Novati et al., 2017; Feng et al., 2023), underwater devices (Zhang et al., 2022; Feng et al., 2024), and nuclear fusion (Degraeve et al., 2022). Furthermore, the adjoint method (Protas, 2008) and PINNs (Mowlavi & Nabi, 2023) are also incorporated in PDE control, but they require an explicit form of the PDE. Currently, the diffusion model used in physical systems’ control (Wei et al., 2024) integrates the learning of entire state trajectories and control sequences, enabling global optimization that incorporates the physical information learned by the model. However, it does not consider the important cases where safety constraints are required.

A.2 SAFE OFFLINE REINFORCEMENT LEARNING

Recently, the offline setting has attracted attention in the field of safe reinforcement learning (RL), as it avoids generating dangerous behaviors through direct interaction with the environment (Achiam et al., 2017; Zhang et al., 2020; Stooke et al., 2020; Liu et al., 2022). CPQ (Xu et al., 2022) is the first practical safe offline RL method that assigns high costs to OOD and unsafe actions and updates the value function as well as the policy only with safe actions. COptiDICE (Lee et al., 2022) is a DICE-based method and corrects the stationary distribution. CDT (Liu et al., 2023b) takes the decision transformer to solve safe offline RL problems as multi-objective optimization. However, these methods lack the ability to model high-dimensional state space. More recent methods like TREBI (Lin et al., 2023) and FISOR (Zheng et al., 2024) solve it through diffusion model planning. But they do not consider the upper bound of the safety score in a probabilistic sense and differ significantly from SafeConPhy in terms of the algorithm.

A.3 CONFORMAL PREDICTION

Conformal prediction (Vovk et al., 2005a) is a statistical framework that constructs prediction intervals guaranteed to contain the true label with a specified probability. Its validity could be compromised, however, by the violation of the core assumption of exchangeability due to distribution shifts in real-world scenarios (Chernozhukov et al., 2018; Hendrycks et al., 2018). Recent studies (Maxime Cauchois & Duchi, 2024) have extended conformal prediction to accommodate various distribution shifts. For example, Tibshirani et al. (2019) proposed weighted conformal prediction to handle covariate shift, where training and test data distributions differ. Podkopaev & Ramdas (2021) introduced reweighted conformal prediction and calibration techniques to address label shift using unlabeled target data. Adaptive conformal inference (Gibbs & Candes, 2021) provides valid prediction sets in online settings with unknown, time-varying distribution shifts without relying on exchangeability. Inspired by previous approaches, SafeConPhy establishes an upper bound on the confidence level by maintaining a weighted score set. Our method ensures the true safety value for a given control sequence lies within the bound without requiring additional assumptions about the model or data distribution, effectively addressing distribution shifts between pre-collected data and the target distribution.

Table 4: **Results of the ablation study.** We compare SafeConPhy with SafeConPhy w/o fine-tuning.

	1D		2D	
	SafeConPhy	w/o fine-tuning	SafeConPhy	w/o fine-tuning
$\mathcal{J} \downarrow$	0.0011	0.0006	-0.7035	-0.6105
$s_{\text{norm}} \downarrow$	0.8388	2.3743	0.6092	0.9096

B ABLATION STUDY

We highlight that one key distinction between our framework and previous safe RL methods lies in the introduction of fine-tuning within Iterative Safety Improvement, which updates the model parameters based on specific control tasks and safety constraints. To further validate the effectiveness of our proposed Iterative Safety Improvement, we conduct experiments using a version of SafeConPhy without the fine-tuning component. As shown in the table, without fine-tuning, SafeConPhy exhibits a significant decline in safety, both in 1D and 2D settings, with the 1D case becoming notably unsafe. This emphasizes the importance and effectiveness of the Iterative Safety Improvement framework in addressing safety control problems.

C LIMITATION AND FUTURE WORK

Firstly, both experiments presented in the paper are not real-world experiments. However, our method is not constrained by specific scenarios, meaning it can be applied to real-world tasks as well, which is our future work. Secondly, we consider extending this method to other generative methods that require constraints, not just the diffusion model. Finally, in the future, we will explore the possibility of developing a stricter bound that sacrifices less accuracy while still ensuring safety.

D VISUALIZATION OF EXPERIMENT RESULTS

D.1 1D BURGERS’ EQUATION

In this section, we provide additional visualizations of the control results for the 1D Burgers’ equation, as shown in Figure 4. In these figures, the top row represents the original trajectories corresponding to the control targets, while the bottom row displays the trajectories controlled by SafeConPhy. It can be observed that SafeConPhy successfully controls the trajectories, preventing boundary violations and guiding them to the desired final state.

D.2 2D INCOMPRESSIBLE FLUID

Here, we provide additional visualizations of the control problems of 2D incompressible fluid. From the figures, we can observe that SafeConPhy can successfully control the smoke to avoid the red hazardous region and reach the target bucket as well.

E ADDITIONAL DETAILS FOR CONFORMAL PREDICTION

Conformal prediction is a flexible framework that provides prediction intervals with guaranteed coverage probabilities for new, unseen data points, under the assumption that the data are exchangeable.

Theoretical Foundations The exchangeability assumption is a cornerstone of conformal prediction. It requires that the order of the data points does not affect their joint distribution, meaning that any permutation of the indices yields an identical distribution. In particular, exchangeability holds for independent and identically distributed (i.i.d.) samples, a common assumption in machine learning tasks. Ensuring exchangeability guarantees the validity of the prediction intervals constructed using conformal prediction.

Implementation Details To implement conformal prediction, the dataset is first split into two subsets: a proper training set (Tr) and a calibration set (Cal). A predictive model μ_θ is trained on the training set using a specified learning algorithm \mathcal{A} . Once trained, the model generates predictions

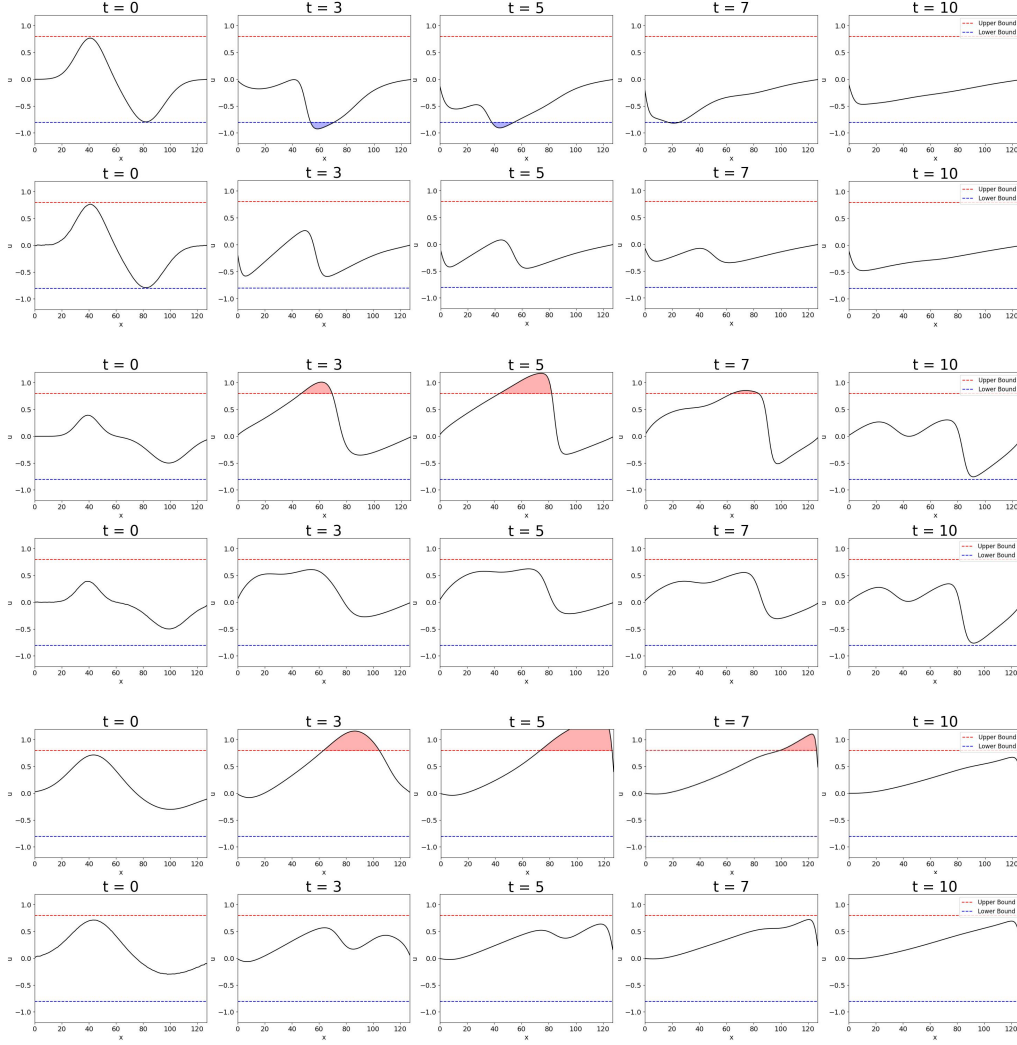


Figure 4: Visualization of the 1D Burgers' equation.

for the calibration set. These predictions are used to compute ‘conformity scores’, which measure the model’s accuracy for each calibration point. Specifically, for each instance i in the calibration set, the conformity score S_i is defined as:

$$S_i = |\mu_\theta(X_i) - Y_i|, \quad i \in \text{Cal}.$$

Additionally, a worst-case score of ∞ is included to account for extreme scenarios.

Then $1 - \alpha$ quantile $q_{1-\alpha}(S)$ of the set of conformity scores is calculated, where α represents the desired significance level (e.g., $\alpha = 0.05$ for 95% confidence interval).

Given a new data point X_{n+1} , the prediction interval for its corresponding output is calculated as:

$$\hat{C}_\alpha(X_{n+1}) = [\mu_\theta(X_{n+1}) - q_{1-\alpha}(S), \mu_\theta(X_{n+1}) + q_{1-\alpha}(S)].$$

This interval provides an estimate for the range within which the true value Y_{n+1} is expected to lie, with a coverage probability of at least $1 - \alpha$. Thus, conformal prediction offers a flexible and robust method for constructing prediction intervals that account for both the model’s accuracy and the variability in the data.

Theoretical Guarantees Conformal prediction provides theoretical guarantees for finite samples (Vovk et al., 2005b; Lei et al., 2016). Specifically, for any new data point, the prediction interval




Figure 5: **Visualizations of the 2D incompressible fluid control problem..**

satisfies the following probabilistic bound:

$$P(Y_{n+1} \in \hat{C}_\alpha(X_{n+1})) \geq 1 - \alpha.$$

This ensures that the true label Y_{n+1} will fall within the predicted interval at least $1 - \alpha$ percent of the time. This framework, based on the assumption of exchangeability, provides a robust method for generating reliable prediction intervals, even in settings with limited sample sizes.

F ADDITIONAL DETAILS FOR 1D EXPERIMENT

F.1 EXPERIMENT SETTING

Following the previous works (Holl et al., 2020; Wei et al., 2024), we generate the 1D Burgers’ equation dataset. During inference, alongside the control sequence $\mathbf{w}(t, x)$, our diffusion model generates states $\mathbf{u}(t, x)$. Our reported evaluation metric \mathcal{J} is always computed by feeding the control $\mathbf{w}(t, x)$ into the ground truth numerical solver to get $\mathbf{u}_{\text{g.t.}}(t, x)$ and computed following Eq. (14). More importantly, we consider the safety constraint and define the safety score as \mathbf{u}^2 . In our experiment, the safety bound is fixed at 0.64, and the details of the 1D Burgers’ equation dataset for safe physical system control problem are listed in Table 5.

Table 5: **Details of 1D Burgers’ equation dataset.**

	Training Set	Calibration Set	Test Set
Unsafe Trajectories	34,985	900	50
Safe Trajectories	4,015	100	0

F.2 MODEL

The model architecture in this experiment follows the Denoising Diffusion Probabilistic Model (DDPM) (Ho et al., 2020). For control tasks, we condition on u_0, u_T and apply guidance to generate the full trajectories of $u_{[0,T]}$, $f_{[0,T]}$ and the safety score s . The hyperparameters for the 2D-Unet architecture are recorded in Table 6.

Table 6: **Hyperparameters of 2D-Unet architecture in 1D experiment.**

Hyperparameter Name	Value
Initial dimension	128
Convolution kernel size	3
Dimension multiplier	[1,2,4,8]
Resnet block groups	1
Attention hidden dimension	32
Attention heads	4
Number of training steps	200000
DDIM sampling iterations	100
η of DDIM sampling	1

G ADDITIONAL DETAILS FOR 2D EXPERIMENT

G.1 EXPERIMENT SETTING

Following works from Holl et al. (2020) and Wei et al. (2024), we use the package `PhiFlow` to generate the 2D incompressible fluid dataset. The control objective and data generation is the same as before (Wei et al., 2024). The main difference between our data and previous ones is that we consider the safety constraint here. We define the safety score as the percentage of smoke passing through a specific region. This reflects the need to limit the amount of pollutants passing through certain areas in real-world scenarios, such as in a watershed.

We simulate the fluid on a 128×128 grid. The selected hazardous region is $[44, 36] \times [40, 64]$. Since the optimal path for smoke, starting from a left-biased position, is likely to pass through this hazardous region, this poses a greater challenge for the algorithm: how to balance safety and achieving a more optimal objective, making this a more difficult problem.

G.2 MODEL

In this paper, the design of the three-dimensional U-net we use is based on the previous work (Ho et al., 2022). In our experiment, we utilize spatio-temporal 3D convolutions. The U-net consists of three key components: a downsampling encoder, a central module, and an upsampling decoder.

The diffusion model conditions on the initial density and uses guidance as previously mentioned to generate the full trajectories of density, velocity, control, the objective \mathcal{J} and the safety score. As \mathcal{J} and s are scalars, we repeat them to match other channels. The hyperparameters for the 3D U-net architecture are listed in Table 7.

Table 7: **Hyperparameters of 3D-Unet architecture in 2D experiments.**

Hyperparameter Name	Value
Number of attention heads	4
Kernel size of conv3d	(3, 3, 3)
Padding of conv3d	(1,1,1)
Stride of conv3d	(1,1,1)
Kernel size of downsampling	(1, 4, 4)
Padding of downsampling	(1, 2, 2)
Stride of downsampling	(0, 1, 1)
Kernel size of upsampling	(1, 4, 4)
Padding of upsampling	(1, 2, 2)
Stride of upsampling	(0, 1, 1)
Number of training steps	200000
DDIM sampling iterations	100
η of DDIM Sampling	1
Intensity of guidance in control	100
Weight of safety term in guidance	10000

H BASELINES

H.1 CDT

Constraints Decision Transformer (CDT) (Liu et al., 2023b) models control as a multi-task regression problem, extending the Decision Transformer (DT) (Chen et al., 2021). It sequentially predicts returns-to-go, costs-to-go, observations, and actions, making actions dependent on previous returns and costs. The authors propose two techniques to adapt the model for safety-constrained scenarios:

1. **Stochastic Policy with Entropy Regularization:** This technique aims to reduce the risk of constraint violations due to out-of-distribution actions. In a deterministic policy, the model selects a single action based on its learned policy, which may result in unsafe actions when faced with states not well represented in the training data. By using a stochastic policy, the model samples actions from a distribution, encouraging the exploration of a wider action space. Entropy regularization further enforces diversity in the sampled actions, making the model more robust in uncertain or underrepresented situations. This approach reduces the likelihood of selecting unsafe actions when faced with states outside the distribution of the training set.
2. **Pareto-Frontier-Based Data Augmentation:** The technique tries to resolve the conflict between maximizing returns and adhering to safety constraints by leveraging a Pareto-frontier of the training data. The Pareto-frontier consists of trajectories that provide the highest possible return under specific safety constraints. Fitting a polynomial to the Pareto-frontier helps identify conflicting high-return and safety constraint pairs, which are then used for augmentation. The augmentation generates synthetic trajectories by relabeling safe trajectories from the Pareto-frontier with higher returns and assigning higher or equal safety constraints. This encourages the model to imitate the most rewarding, safe trajectories when the desired return given the safety constraint is infeasible.

In the 2D incompressible fluid experiment, the model fails to extrapolate safe trajectories with higher returns due to the training data **covering a broad range of costs, while the desired cost lies in a narrow range**. The augmentation treats all safe trajectories on the Pareto-frontier equally, without emphasizing the region of interest. To investigate, we filtered the training data to include only safe trajectories under the desired safety bound and retrained the model, showing that the data complexity exceeds CDT’s capacity. This complexity, however, could potentially be addressed by SafeConPhy, which is indicated in Table 3. We use the official CDT implementation and follow DT guidelines to sweep desired returns and cost constraints in testing time. In the 1D Burgers experiment, we modify the control objective from the original mean squared error \mathcal{J} between the prediction and target, to an exponential form $\exp(-\mathcal{J})$. This new objective is bounded within $[0, 1]$, which better aligns with the reward-maximizing setup in reinforcement learning used by CDT. For the 2D incompressible fluid setup, the state, action, and cost prediction heads each consist of 3-layer MLPs with the transformer’s hidden dimension as the inner size.

Table 8: **Hyperparameters of 1D CDT.**

1D Burgers’	All data	Safe filtered
State Dimension	256	256
Action Dimension	128	128
Hidden Dimension	1024	1024
Number of Transformer Blocks	2	2
Number of Attention Heads	8	8
Horizon (Sequence Length)	5	5
Learning Rate	1e-4	1e-4
Batch Size	64	64
Weight Decay	1e-5	1e-5
Learning Steps	1,000,000	1,000,000
Learning Rate Warmup Steps	500	500
Pareto-Frontier Fitted Polynomial Degree	0	4
Augmentation Data Percentage	0.3	0.3
Max Augment Reward	10.0	10.0
Min Augment Reward	1.0	1.0
Target Entropy	-128	-128
Testing Time Sweep Returns	9.0, 9.9	9.0, 9.9
Testing Time Sweep Costs	0.0, 1.0, 2.0, 3.0	0.0, 1.0, 2.0, 3.0

Table 9: **Hyperparameters of 2D CDT.**

2D incompressible fluid	All data	Safe filtered
State Dimension	$3 \times 64 \times 64$	$3 \times 64 \times 64$
Action Dimension	$2 \times 64 \times 64$	$2 \times 64 \times 64$
Hidden Dimension	512	512
Number of Transformer Blocks	3	3
Number of Attention Heads	8	8
Horizon (Sequence Length)	10	10
Learning Rate	1e-4	1e-4
Batch Size	8	8
Weight Decay	1e-5	1e-5
Learning Steps	1,000,000	1,000,000
Learning Rate Warmup Steps	500	500
Pareto-Frontier Fitted Polynomial Degree	4	4
Augmentation Data Percentage	0.3	0.3
Max Augment Reward	32.0	32.0
Min Augment Reward	1.0	1.0
Target Entropy	$-(2 \times 64 \times 64)$	$-(2 \times 64 \times 64)$
Testing Time Sweep Returns	18.0, 32.0	18.0, 32.0
Testing Time Sweep Costs	0.0, 0.1, 0.2	0.0, 0.1, 0.2

H.2 BC-ALL

The Behavior Cloning (BC) algorithm, introduced by (Pomerleau, 1988), is a foundational technique in imitation learning. BC is designed to derive policies directly from expert demonstrations, utilizing supervised learning to associate states with corresponding actions. This method eliminates the necessity for exploratory steps commonly required in reinforcement learning by replicating the actions observed in expert demonstrations. One of the significant advantages of BC is that it does not involve interacting with the environment during the training phase, which streamlines the learning process and diminishes the demand for computational resources.

In this approach, a policy network is trained using standard supervised learning strategies aimed at reducing the discrepancy between the actions predicted by the model and those performed by the expert in the dataset. The commonly used loss function for this purpose is the mean squared error between the predicted actions and expert actions. The dataset for training comprises state-action pairs harvested from these expert demonstrations. In the work, we employ the implementation as Liu et al. (2023a).

H.3 BC-SAFE

Following the baseline in Liu et al. (2023a), the BC-Safe is fed with only safe trajectories filtered from the training dataset, satisfies most safety requirements, although with conservative performance and lower rewards. Others are same as BC-All, except for the safe trajectories.

H.4 SL-LAG

Hwang et al. (2022) introduces a supervised learning (SL) based method to control PDE systems. It first trains a neural surrogate model to capture the PDE dynamics, which includes a VAE to compress PDE states and controls into the latent space and another model to learn the PDE’s time evolution in the latent space. To obtain the optimal control sequence, SL can compute the gradient $\nabla_{\mathbf{w}} \mathcal{J}$, where \mathcal{J} is the control objective and f is the input control sequence. Then iterative gradient optimization can be executed to improve the control sequence.

To ensure that the optimal control is compatible with the hard constraint in our experiments, we follow Chow et al. (2018) to apply the Lagrange optimization method to the constrained optimization. Specifically, we iteratively solve the optimization problem below:

$$\max_{\lambda \geq 0} \min_{\mathbf{w}} \mathcal{J}(\mathbf{w}) + \lambda(s(\mathbf{u}(\mathbf{w})) - c). \quad (16)$$

We denote the modified SL method SL-Lag.

Table 10: **Hyperparameters of network architecture and training for SL-Lag in 1D Burgers’ experiment.**

Hyperparameter name	Value
Initialization value of \mathbf{w}	0.001
Optimizer of \mathbf{w}	LBFGS
Learning rate of \mathbf{w}	0.1
Initialization value of the Lagrange multiplier λ	0
Optimizer of \mathbf{w}	plain GD
Learning rate of λ	10
Iteration of λ	2
Loss function	MSE

H.5 TREBI

In Lin et al. (2023), the diffusion model is adopted for the planning task under safety budgets. It generates trajectory under safety constraints using classifier-guidance Dhariwal & Nichol (2021) by adding a safety loss to the reward guidance following Diffuser Janner et al. (2022).

However, the original setting is different from our experiments. In our 1D Burgers’ equation control, our objective is that a certain state equals the target state, which in-painting diffusion condition Janner et al. (2022) is more appropriate for. Furthermore, as in our method and in (Wei et al., 2024), a conditional diffusion model can be learned to tackle the objective more directly. In addition, TREBI follows the setting in Diffuser where the interaction with the environment is allowed which in our experiments becomes an MPC method. Note that the reported results of our method do not involve interaction with the surrogate model (though our method can easily adapted to be an MPC method). Thus, the results of TREBI in Table 2 and 3 have an unfair advantage.

Therefore, for 1D Burgers’ experiment, we conducted different experiments on TREBI including (1) planning multiple times with interaction with the surrogate model or (2) planning only once, and with target state conditioning or target state guidance. The target state guidance + planning multiple times turned out the best and is reported in Table 2. For 2D smoke control, the target is not a state constraint but a reward, and the planning multiple-step setting is too computationally expensive. To this end, we use reward guidance with planning one single time, which is identical to the ablation study of our method in Table 4. The hyperparameters of the 1D experiment are reported in Table 11, and those of 2D are the same as reported before.

Table 11: **Hyperparameters of network architecture and training for TREBI in 1D Burgers’ experiment.**

Hyperparameter name	Value
Reward guidance intensity	50
Safety guidance intensity	1
Cost budget	0.64
Number of guidance steps	10
Denoising steps	200
Sampling algorithm	DDPM
U-Net dimension	64
U-Net dimension mltiplications	1, 2, 4, 8
Planning horizon	8 steps
Optimizer	Adam
Learning rate	0.0002
Batch size	16
Loss function	MSE

H.6 PID

Propercentageal Integral Derivative (PID) control (Li et al., 2006) is a versatile and effective method widely employed in numerous control scenarios. For 1D control task, we mainly implement the PID baseline adapted from Wei et al. (2024). More detailed configurations can be found in Table 12

Table 12: **Hyperparameters of network architecture and training for ANN PID.**

Hyperparameter name	Value
Kernel size of conv1d	3
Padding of conv1d	1
Stride of conv1d	1
Activation function	Softsign
Batch size	16
Optimizer	Adam
Learning rate	0.0001
Loss function	MAE